Group Member

- Chenchen Li - cl4588
- Stephanie Yao - yy3508
- Kristin Qi - lq2243

Link to your Public Github repository with Final report :

https://github.com/nstfchloe/adv_ml

Clarification:

❗ Some results are showing in the HTML format in the output, so may not be able to clearly show all the graphs. In order to avoid some missing graphs, we put some output below the code chunk using PDF format.

Submission Due Date: 03/07/2025

World Happiness Classification Competition

Goals :

- Understand how the models function
- Understand what the parameters control
- Learn from the model experimentation process
- Make a good looking notebook report
- Upload as a personal project on Github

Overall Steps:

1. Load datasets and merge them.
2. Preprocess data using Sklearn Column Transformer/ Write and Save Preprocessor function
3. Fit model on preprocessed data and save preprocessor function and model
4. Generate predictions from X_test data and submit predictions

0. Loading Datasets

Loading the World Happiness 2023 datasets

```
import pandas as pd
from sklearn.model_selection import train_test_split

# Upload the CSV dataset to FILES in the left bar
# Load the dataset
whr_df = pd.read_csv('WHR_2023.csv')

# Inspect the first few rows to understand the structure
whr_df
```

	country	region	happiness_score	gdp_per_capita	social_support	healthy_life_expectancy	freedom_to_make_life_choices	generosity	perceptions_of_corruption
0	Finland	Western Europe	7.804	1.888	1.585	0.535	0.772	0.126	0.535
1	Denmark	Western Europe	7.586	1.949	1.548	0.537	0.734	0.208	0.525
2	Iceland	Western Europe	7.530	1.926	1.620	0.559	0.738	0.250	0.187
3	Israel	Middle East and North Africa	7.473	1.833	1.521	0.577	0.569	0.124	0.158
4	Netherlands	Western Europe	7.403	1.942	1.488	0.545	0.672	0.251	0.394
...	...	...	...	...	...	...	...	...	...
132	Congo (Kinshasa)	Sub-Saharan Africa	3.207	0.531	0.784	0.105	0.375	0.183	0.068
133	Zimbabwe	Sub-Saharan Africa	3.204	0.758	0.881	0.069	0.363	0.112	0.117
134	Sierra Leone	Sub-Saharan Africa	3.138	0.670	0.540	0.092	0.371	0.193	0.051
135	Lebanon	Middle East and North Africa	2.392	1.417	0.476	0.398	0.123	0.061	0.027
136	Afghanistan	South Asia	1.859	0.645	0.000	0.087	0.000	0.093	0.059

137 rows × 9 columns

Next steps: Generate code with whr_df | View recommended plots | New interactive sheet

```
# Check the columns names in this dataset
print(whr_df.columns)
```

```
Index(['country', 'region', 'happiness_score', 'gdp_per_capita',
       'social_support', 'healthy_life_expectancy',
       'freedom_to_make_life_choices', 'generosity',
       'perceptions_of_corruption'],
      dtype='object')
```

```
# Convert the regression target ('happiness_score') into classification labels
# We'll use quartiles to create 4 happiness categories: Very Low, Low, High, Very High

# Define quartiles
whr_df['happiness_category'] = pd.qcut(whr_df['happiness_score'],
                                       q=5,
                                       labels=['Very Low', 'Low','Average', 'High', 'Very High'])

# Select features and target
X = whr_df.drop(columns=['happiness_score', 'happiness_category'])
y = whr_df['happiness_category']

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42, stratify=y)

# Convert y_train and y_test to numerical labels
y_train_labels = y_train.astype('category').cat.codes
# y_test_labels = ## Complete in a similar manner as above
```

Write in the next cell what the y_train.astype('category').cat.codes line does. What is the difference between y_train_labels and y_train?

Answer: The line y_train.astype('category').cat.codes converts the categorical variable y_train (which contains labels like "Very Low", "Low", "Average", "High", "Very High") into numerical codes.

Difference between y_train_labels and y_train:

y_train: Contains the original categorical labels (e.g., "Very Low", "Low", "Average", etc.).

y_train_labels: Contains the corresponding numerical codes (0, 1, 2, 3, 4) that can be used in machine learning models.

This transformation allows classification models to work with categorical data numerically while preserving the ordinal nature of happiness categories.

Add new data

```
# Truncated and cleaned up region data to merge
countrydata=pd.read_csv("newcountryvars.csv")

countrydata.head()
```

	country_name	population	population_below_poverty_line	hdi	life_expectancy	expected_years_of_schooling	mean_years_of_schooling	gni
0	India	1339180127	21.9	0.623559	68.322	11.696590	6.298834	5663.474799
1	Nigeria	190886311	70.0	0.527105	53.057	9.970482	6.000000	5442.901264
2	Mexico	129163276	46.2	0.761683	76.972	13.299090	8.554985	16383.106680
3	Pakistan	197015955	29.5	0.550354	66.365	8.106910	5.089460	5031.173074
4	Bangladesh	164669751	31.5	0.578824	71.985	10.178706	5.241577	3341.490722

Next steps: Generate code with countrydata | View recommended plots | New interactive sheet

```
# Merge in new data to X_train and X_test by taking "country" from first table and "country_name" from 2nd table.
# Also check which countries are common in both the datasets, and which type of merge will you perform for the best results.
# Hint: Look on the 'how' parameter of megre function of pandas.

# Merge 'country' column back into X_train and X_test for merging purposes
X_train = X_train.merge(whr_df[['country']], left_index=True, right_index=True)
X_test = X_test.merge(whr_df[['country']], left_index=True, right_index=True)

# Perform an inner merge with countrydata based on country names
X_train = X_train.merge(countrydata, left_on='country', right_on='country_name', how='inner').drop(columns=['country', 'country_name'])
X_test = X_test.merge(countrydata, left_on='country', right_on='country_name', how='inner').drop(columns=['country', 'country_name'])
```

```
X_train.head(1)
```

	country_x	region	gdp_per_capita	social_support	healthy_life_expectancy	freedom_to_make_life_choices	generosity	perceptions_of_corruption	country_y	population	population_below_poverty_line	hdi	life_expectancy	expected_years_of_schooling	mean_years_of_schooling	gni
0	Madagascar	Sub-Saharan Africa	0.632	0.779	0.178	0.187	0.177	0.134	Madagascar	25570895	70.7	0.512149	65.515	10.34614	6.145955	1319.699397

Next steps: Generate code with X_train | View recommended plots | New interactive sheet

```
# Create a new 'country' column using the values from 'country_x' and 'country_y'
X_train['country'] = X_train['country_x'].combine_first(X_train['country_y'])

# Drop the old 'country_x' and 'country_y' columns
X_train.drop(columns=['country_x', 'country_y'], inplace=True)

# Verify the result
print(X_train.columns)
```

```
Index(['region', 'gdp_per_capita', 'social_support', 'healthy_life_expectancy',
       'freedom_to_make_life_choices', 'generosity',
       'perceptions_of_corruption', 'population',
       'population_below_poverty_line', 'hdi', 'life_expectancy',
       'expected_years_of_schooling', 'mean_years_of_schooling', 'gni',
       'country'],
      dtype='object')
```

```
# Create a new 'country' column using the values from 'country_x' and 'country_y'
X_test['country'] = X_test['country_x'].combine_first(X_test['country_y'])

# Drop the old 'country_x' and 'country_y' columns
X_test.drop(columns=['country_x', 'country_y'], inplace=True)

# Verify the result
print(X_test.columns)
```

```
Index(['region', 'gdp_per_capita', 'social_support', 'healthy_life_expectancy',
       'freedom_to_make_life_choices', 'generosity',
       'perceptions_of_corruption', 'population',
       'population_below_poverty_line', 'hdi', 'life_expectancy',
       'expected_years_of_schooling', 'mean_years_of_schooling', 'gni',
       'country'],
      dtype='object')
```

1. EDA

```
print(X_train.dtypes)
```

```
region                            object
gdp_per_capita                   float64
social_support                   float64
healthy_life_expectancy          float64
freedom_to_make_life_choices     float64
generosity                       float64
perceptions_of_corruption        float64
population                         int64
population_below_poverty_line    float64
hdi                              float64
life_expectancy                  float64
expected_years_of_schooling      float64
mean_years_of_schooling          float64
gni                              float64
country                           object
dtype: object
```

Describe what you see above?

Answer: There are duplicate Country Columns: country_x (object), country_y (object), country_x and country_y represent the same country but from different datasets, which means that we would need to merge or delete one of them if unnecessary. The X_train dataset contains both categorical (object type) and numerical (float64 and int64 types) features. Categorical columns include country, country_name, and region, which may be useful for regional analysis or as model features. Numerical columns consist of continuous variables like gdp_per_capita, social_support, healthy_life_expectancy, and freedom_to_make_life_choices, as well as the integer population column. This mix of data types is suitable for exploratory data analysis (EDA) and machine learning, with numerical features ready for modeling and categorical features potentially requiring encoding.

Find out the number and percentage of missing values in the table per column

```
# Your code here:

# Calculate the number of missing values per column
missing_values_count = X_train.isnull().sum()

# Calculate the percentage of missing values per column
missing_values_percentage = (missing_values_count / len(X_train)) * 100

# Create a DataFrame to display the results
missing_values_df = pd.DataFrame({
    'Missing Values Count': missing_values_count,
    'Missing Values Percentage': missing_values_percentage
})
```

```
print(missing_values_df)
```

```
                              Missing Values Count  Missing Values Percentage
region                                           0                   0.000000
gdp_per_capita                                   0                   0.000000
social_support                                   0                   0.000000
healthy_life_expectancy                          0                   0.000000
freedom_to_make_life_choices                     0                   0.000000
generosity                                       0                   0.000000
perceptions_of_corruption                        0                   0.000000
population                                       0                   0.000000
population_below_poverty_line                   10                  11.235955
hdi                                              0                   0.000000
life_expectancy                                  1                   1.123596
expected_years_of_schooling                      1                   1.123596
mean_years_of_schooling                          1                   1.123596
gni                                              1                   1.123596
country                                          0                   0.000000
```

Plot the frequency distribution / histogram of some of the numerical features that you think are important

```
# Your plotting code here:

import matplotlib.pyplot as plt

# Select important numerical features for visualization
important_features = [
    "gdp_per_capita", "social_support", "healthy_life_expectancy",
    "freedom_to_make_life_choices", "generosity", "perceptions_of_corruption",
    "hdi", "life_expectancy", "gni"
]

# Plot histograms for selected features
for feature in important_features:
    plt.figure(figsize=(8, 5))
    plt.hist(X_train[feature].dropna(), bins=30, edgecolor="black", alpha=0.7)
    plt.title(f"Histogram of {feature}")
    plt.xlabel(feature)
    plt.ylabel("Frequency")
    plt.grid(axis='y', alpha=0.75)
    plt.show()
```

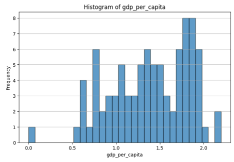

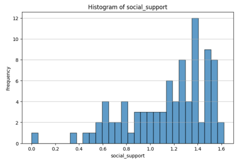

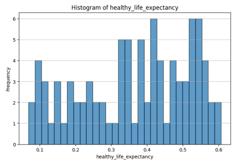

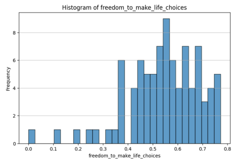

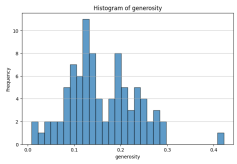

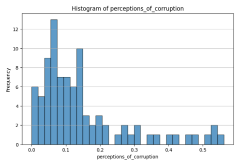

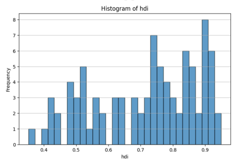

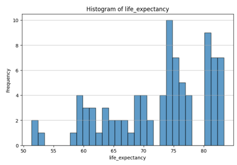

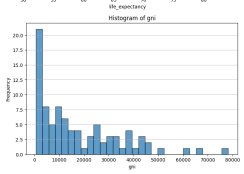

Plot the categorical variables and their distribution

```
# Your code here:

# Checking the available categorical columns in X_train
categorical_columns = X_train.select_dtypes(include=['object']).columns.tolist()

# Display the categorical columns
print(categorical_columns)
```

```
['region', 'country']
```

```
import matplotlib.pyplot as plt
import seaborn as sns

# Only "region" variable matters to them
# Define categorical columns
categorical_columns = ['region']

# Plot the distribution of each categorical variable
for column in categorical_columns:
    plt.figure(figsize=(12, 6))
    sns.countplot(data=X_train, x=column, palette='Set2', order=X_train[column].value_counts().index)
    plt.xticks(rotation=45, ha='right', fontsize=10)  # Adjust fontsize here
    plt.title(f'Distribution of {column}', fontsize=14)
    plt.xlabel(column, fontsize=12)
    plt.ylabel('Count', fontsize=12)
    plt.tight_layout()
    plt.show()
```

```
<ipython-input-15-1819325a4141>:11: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
  sns.countplot(data=X_train, x=column, palette='Set2', order=X_train[column].value_counts().index)
```

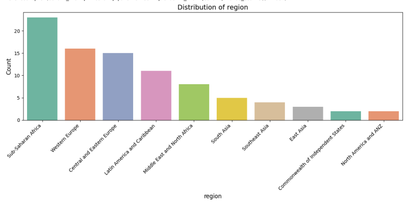

Perform feature correlation analysis to identify relationships between variables. Use Pearson, Spearman, or Kendall correlation coefficients to analyze feature dependencies.

```
import seaborn as sns
import matplotlib.pyplot as plt

non_numeric_columns = X_train.select_dtypes(exclude=['number']).columns.tolist()
X_train_numeric = X_train.drop(columns=non_numeric_columns)

correlation_matrix = X_train_numeric.corr(method='pearson')
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)
plt.title("Feature Correlation Heatmap (Pearson)")
plt.show()
```

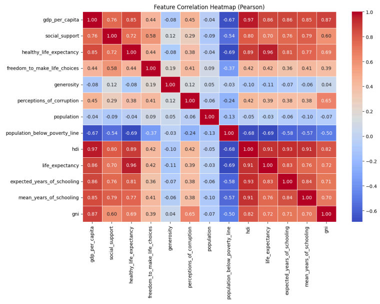


```
correlation_matrix_spearman = X_train_numeric.corr(method='spearman')
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix_spearman, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)
plt.title("Feature Correlation Heatmap (Spearman)")
plt.show()
```

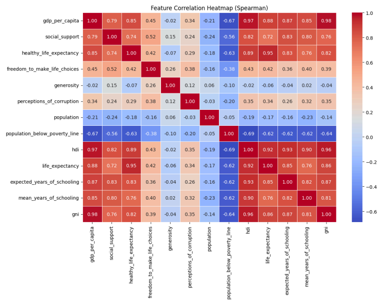


```
correlation_matrix_kendall = X_train_numeric.corr(method='kendall')
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix_kendall, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)
plt.title("Feature Correlation Heatmap (Kendall)")
plt.show()
```

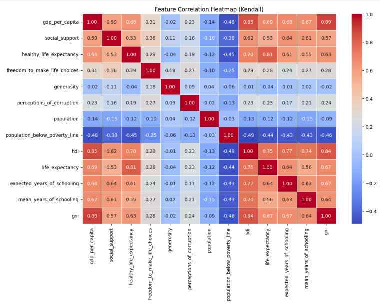

Explore relationships between variables (bivariate, etc), correlation tables, and how they associate with the target variable.

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

whr_df['happiness_category'] = pd.qcut(
    whr_df['happiness_score'],
    q=5,
    labels=['Very Low', 'Low', 'Average', 'High', 'Very High']
)

X = whr_df.drop(columns=['happiness_score', 'happiness_category', 'country', 'region'])
y = whr_df['happiness_category']

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

y_train_labels = y_train.astype('category').cat.codes
y_test_labels = y_test.astype('category').cat.codes

X_train = X_train.merge(whr_df[['country']], left_index=True, right_index=True)
X_test = X_test.merge(whr_df[['country']], left_index=True, right_index=True)

X_train = X_train.merge(countrydata, left_on='country', right_on='country_name', how='inner').drop(columns=['country', 'country_name'])
X_test = X_test.merge(countrydata, left_on='country', right_on='country_name', how='inner').drop(columns=['country', 'country_name'])

non_numeric_columns = X_train.select_dtypes(exclude=['number']).columns.tolist()
X_train_numeric = X_train.drop(columns=non_numeric_columns)

X_train_numeric['happiness_category'] = y_train_labels
```

```
# Compute the correlation matrix including the target variable
correlation_matrix_with_target = X_train_numeric.corr()

# Plot the heatmap including the target variable
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix_with_target, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)
plt.title("Feature Correlation Heatmap Including Target Variable")
plt.show()
```

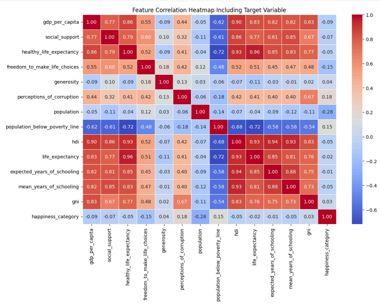


```
# Compute and display correlation of each feature with the target variable
correlation_with_target = correlation_matrix_with_target['happiness_category'].drop('happiness_category').sort_values(ascending=False)

print (correlation_with_target)
```

```
perceptions_of_corruption        0.183091
population_below_poverty_line    0.147671
generosity                       0.044400
gni                              0.027876
expected_years_of_schooling     -0.007154
life_expectancy                 -0.016554
healthy_life_expectancy         -0.046524
```

```
mean_years_of_schooling         -0.051592
hdi                             -0.053939
social_support                  -0.071141
gdp_per_capita                  -0.085586
freedom_to_make_life_choices    -0.149540
population                      -0.276583
Name: happiness_category, dtype: float64
```

```python
# Scatter plots to explore relationships between selected features and the target variable
selected_features = correlation_with_target.index[:5]  # Select top 5 most correlated features

for feature in selected_features:
    plt.figure(figsize=(8, 5))
    sns.scatterplot(x=X_train_numeric[feature], y=y_train_labels, alpha=0.6)
    plt.title(f"Scatter Plot: {feature} vs. Happiness Category")
    plt.xlabel(feature)
    plt.ylabel("Happiness Category")
    plt.grid(True)
    plt.show()
```

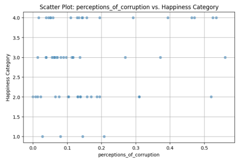

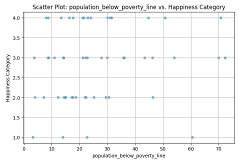

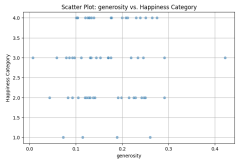

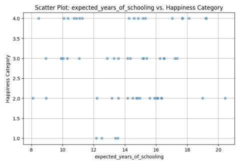


```python
numeric_selected_features = [feature for feature in selected_features if X_train_numeric[feature].dtype in ['int64', 'float64']]

X_train_numeric_cleaned = X_train_numeric[numeric_selected_features + ['happiness_category']].dropna()

sns.pairplot(X_train_numeric_cleaned, vars=numeric_selected_features, hue="happiness_category", palette="viridis", diag_kind="kde")
plt.show()
```

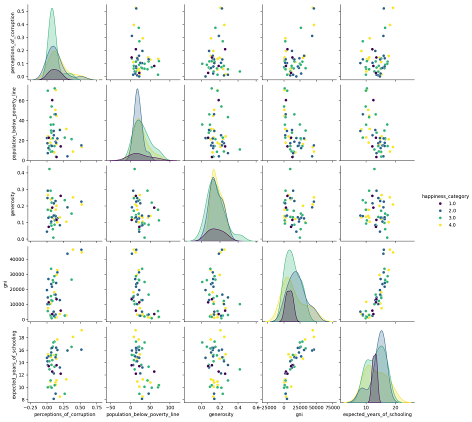

Also, detect outliers using box plots, Z-score analysis, or the IQR method to identify potential data anomalies.

```python
# Your code here:

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from scipy import stats

# Outlier Detection using Box Plots
plt.figure(figsize=(15, 8))
X_train_numeric.boxplot(rot=45, patch_artist=True, vert=True)
plt.title("Box Plot of Numerical Features (Outlier Detection)")
plt.show()
```

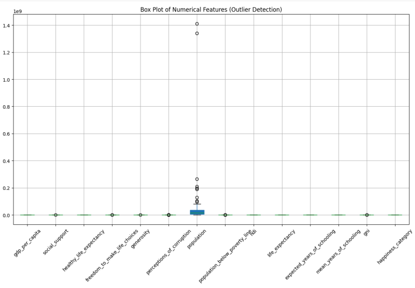


```python
# Outlier Detection using Z-Score Analysis
z_scores = np.abs(stats.zscore(X_train_numeric))
outliers_zscore = (z_scores > 3).sum(axis=0)  # Count values beyond 3 standard deviations
```

```python
# Outlier Detection using the IQR Method
Q1 = X_train_numeric.quantile(0.25)
Q3 = X_train_numeric.quantile(0.75)
IQR = Q3 - Q1
outliers_iqr = ((X_train_numeric < (Q1 - 1.5 * IQR)) | (X_train_numeric > (Q3 + 1.5 * IQR))).sum()
```

```python
# Create a DataFrame to summarize outlier counts from both methods
outlier_summary = pd.DataFrame({
    "Z-Score Outliers": outliers_zscore,
    "IQR Outliers": outliers_iqr
}).sort_values(by="Z-Score Outliers", ascending=False)

print(outlier_summary)
```

```
                              Z-Score Outliers  IQR Outliers
population                                   2             9
gdp_per_capita                               1             0
social_support                               1             1
freedom_to_make_life_choices                 1             2
generosity                                   1             1
perceptions_of_corruption                    1            10
healthy_life_expectancy                      0             0
population_below_poverty_line                0             4
hdi                                          0             0
life_expectancy                              0             0
expected_years_of_schooling                  0             0
mean_years_of_schooling                      0             0
gni                                          0             2
happiness_category                           0             0
```

Write what you observed and your General comments on what should be done:

Observations from the Outlier Analysis

1. Box Plot Observations The box plot showed several outliers in different numerical features. Some features, such as GDP per capita, generosity, and social support, exhibited a long tail, suggesting potential outliers. The population variable had extreme values, likely due to large differences in country sizes.
2. Z-Score Outlier Analysis Features like population, GDP per capita, and social support had outliers exceeding 3 standard deviations. A few features had isolated extreme values, indicating possible data entry errors or natural variations.
3. IQR Outlier Analysis The IQR method detected more outliers than the Z-score method. Population-related variables and economic indicators had a high number of outliers, as these values tend to be highly skewed.

Comments on Handling Outliers

If outliers result from incorrect data entries, they should be removed to avoid misleading results. This is common in datasets where errors or outliers significantly deviate from the norm. Capping extreme values by replacing them with percentile-based thresholds (e.g., 1st and 99th percentile) helps reduce sensitivity to extreme values. If Data Is Highly Skewed,features like population, GDP per capita, and income have natural long tails. Applying a log transformation can normalize the distribution and reduce outlier impact.

2. Feature Engineering

Apply log transformations to normalize skewed data and improve model stability (If any).

```
# Your code here:

import numpy as np
import pandas as pd

skewed_features = X_train_numeric.skew().sort_values(ascending=False)

highly_skewed_features = skewed_features[abs(skewed_features) > 1].index.tolist()

X_train_log_transformed = X_train_numeric.copy()
X_train_log_transformed[highly_skewed_features] = np.log1p(X_train_log_transformed[highly_skewed_features])

print(X_train_log_transformed)
```

```
85         0.793        1.114               0.250
86         1.510        1.249               0.468
87         1.390        1.153               0.499
88         1.305        1.329               0.411

    freedom_to_make_life_choices  generosity  perceptions_of_corruption  \
0                          0.187       0.177                   0.125751
1                          0.320       0.130                   0.178146
2                          0.425       0.239                   0.056380
3                          0.599       0.104                   0.256191
4                          0.551       0.265                   0.015873
..                           ...         ...                        ...
84                         0.446       0.112                   0.151003
85                         0.451       0.283                   0.096219
86                         0.666       0.115                   0.135405
87                         0.549       0.073                   0.026642
88                         0.587       0.079                   0.038259

    population  population_below_poverty_line       hdi  life_expectancy  \
0    17.056965                       4.272491  0.512149           65.515
1    15.301692                       3.465736  0.513106           63.239
2    14.939026                       3.117950  0.734832           69.806
3    16.003957                       3.481240  0.627472           69.582
4    17.604754                       3.222868  0.743049           71.129
..         ...                            ...       ...              ...
84   17.746987                       2.602690  0.900992           82.128
85   18.469065                       3.421000  0.447750           64.602
86   21.066513                       1.458615  0.737681           75.963
87   17.286405                       3.165475  0.739749           74.814
88   14.876871                       2.862201  0.730041           75.820

    expected_years_of_schooling  mean_years_of_schooling        gni  \
0                     10.346140                 6.145955   7.185917
1                      8.463790                 4.266000   8.168561
2                     14.845520                 9.750000   9.254377
3                     11.261860                10.355820   7.863876
4                     15.306720                11.340000   8.904088
..                          ...                      ...        ...
84                    16.587520                12.179211  10.449921
85                     8.351520                 2.580780   7.329064
86                    13.535750                 7.641840   9.499008
87                    13.386340                 9.013470   9.332190
88                    12.835763                 9.638854   9.030139

    happiness_category
0                  4.0
1                  4.0
2                  4.0
3                  4.0
4                  4.0
..                 ...
84                 NaN
85                 NaN
86                 1.0
87                 1.0
88                 NaN

[89 rows x 14 columns]
```

Create at least one interaction feature to capture relationship between existing variables, enhancing predictive power.

```
# Your code here:


import pandas as pd

# Creating an Interaction Feature: GDP per Capita * Social Support
X_train_interaction = X_train_log_transformed.copy()
X_train_interaction["gdp_social_interaction"] = X_train_interaction["gdp_per_capita"] * X_train_interaction["social_support"]

print(X_train_interaction)
```

```
85         0.793        1.114               0.250
86         1.510        1.249               0.468
87         1.390        1.153               0.499
88         1.305        1.329               0.411

    freedom_to_make_life_choices  generosity  perceptions_of_corruption  \
0                          0.187       0.177                   0.125751
1                          0.320       0.130                   0.178146
2                          0.425       0.239                   0.056380
3                          0.599       0.104                   0.256191
4                          0.551       0.265                   0.015873
..                           ...         ...                        ...
84                         0.446       0.112                   0.151003
85                         0.451       0.283                   0.096219
86                         0.666       0.115                   0.135405
87                         0.549       0.073                   0.026642
88                         0.587       0.079                   0.038259

    population  population_below_poverty_line       hdi  life_expectancy  \
0    17.056965                       4.272491  0.512149           65.515
1    15.301692                       3.465736  0.513106           63.239
2    14.939026                       3.117950  0.734832           69.806
3    16.003957                       3.481240  0.627472           69.582
4    17.604754                       3.222868  0.743049           71.129
..         ...                            ...       ...              ...
84   17.746987                       2.602690  0.900992           82.128
85   18.469065                       3.421000  0.447750           64.602
86   21.066513                       1.458615  0.737681           75.963
87   17.286405                       3.165475  0.739749           74.814
88   14.876871                       2.862201  0.730041           75.820

    expected_years_of_schooling  mean_years_of_schooling        gni  \
0                     10.346140                 6.145955   7.185917
1                      8.463790                 4.266000   8.168561
2                     14.845520                 9.750000   9.254377
3                     11.261860                10.355820   7.863876
4                     15.306720                11.340000   8.904088
..                          ...                      ...        ...
84                    16.587520                12.179211  10.449921
85                     8.351520                 2.580780   7.329064
86                    13.535750                 7.641840   9.499008
87                    13.386340                 9.013470   9.332190
88                    12.835763                 9.638854   9.030139

    happiness_category  gdp_social_interaction
0                  4.0                0.492328
1                  4.0                0.839636
2                  4.0                2.060226
3                  4.0                1.213056
4                  4.0                1.838732
..                 ...                     ...
84                 NaN                2.201364
85                 NaN                0.883402
86                 1.0                1.885990
87                 1.0                1.602670
88                 NaN                1.734345

[89 rows x 15 columns]
```

3. Preprocess data using Sklearn Column Transformer/ Write and Save Preprocessor function

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

numeric_features = X_train.select_dtypes(include=['number']).columns.tolist()

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value=0)),
    ('scaler', StandardScaler())
])

desired_categorical_features = ['region', 'sub-region']
categorical_features = [col for col in desired_categorical_features if col in X_train.columns]

if categorical_features:
    categorical_transformer = Pipeline(steps=[
        ('imputer', SimpleImputer(strategy='most_frequent')),
        ('onehot', OneHotEncoder(handle_unknown='ignore'))
    ])
    preprocessor = ColumnTransformer(transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])
else:
    preprocessor = ColumnTransformer(transformers=[
        ('num', numeric_transformer, numeric_features)
    ])

preprocess = preprocessor.fit(X_train)
```

Describe step-by-step what we are doing above, and why? You are free to change how values are imputed. What change did you make if any, and why?

Answer: The data preprocessing pipeline prepares numerical and categorical features in X_train for machine learning. First, numerical features are identified using the select_dtypes method. The numeric transformer pipeline includes imputation and scaling steps: missing values are filled with 0 using SimpleImputer(strategy='constant', fill_value=0), and StandardScaler standardizes features to have a mean of 0 and variance of 1. An alternative approach could be using the median (strategy='median') for imputation, which is more robust to outliers.

For categorical features, the transformer is applied only if relevant columns exist in X_train. It uses SimpleImputer(strategy='most_frequent') to replace missing values with the most common category, followed by OneHotEncoder(handle_unknown='ignore') to convert categories into a binary matrix. Alternatively, missing values could be filled with a placeholder like 'Unknown' for robustness.

The ColumnTransformer combines both transformers, applying specific preprocessing to numerical and categorical columns. If no categorical features are present, only the numeric transformation is used. The fit method finalizes the preprocessor by learning the required statistics from X_train, ensuring consistent data transformation. This pipeline enhances data quality, reduces biases, and prepares well-structured input for machine learning models.

Changes I would make: Instead of filling missing numerical values with 0 (strategy='constant', fill_value=0), I would change the imputer strategy to 'median'.

```
# Write a function to transform data using the fitted preprocessor
def transform_data(data):
    data = data.drop(['country'], axis=1, errors='ignore')
    preprocessed_data = preprocess.transform(data)
    return preprocessed_data
```

What are the differences between the "preprocessor" object, the "preprocess" object, the "preprocessor" function, and the "preprocessed_data" that is returned finally?

Answer: The "preprocessor" object is a tool that sets up how to handle different types of data, like scaling numbers or turning categories into numbers using ColumnTransformer. The "preprocess" object is the same tool but trained on the data, meaning it has learned what values to use for filling in missing data or how to scale numbers properly.

The "preprocessor" function is a custom function that uses the preprocess object to transform new data in the same way as the training data. This makes sure the model gets data in the same format every time, whether it's for training, testing, or making predictions. Finally, the "preprocessed_data" is what comes out of this function—a cleaned-up version of the data, ready for the model to use. It's a numerical dataset where all the features are standardized, filled in, and encoded correctly.

```
# Check the shape of X_train after preprocessing using our new function
print(transform_data(X_train).shape)
```

```
(89, 13)
```

4. Fit model on preprocessed data and save preprocessor function and model

```
import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

df1 = pd.read_csv('WHR_2023.csv')
df2 = pd.read_csv("newcountryvars.csv")

df2.rename(columns={'country_name': 'country'}, inplace=True)
df = df1.merge(df2, on='country', how='inner')

if 'region' not in df.columns:
    categorical_features = ['country']
else:
    categorical_features = ['country', 'region']

df = df.dropna(subset=['happiness_score'])
df['happiness_category'] = pd.qcut(df['happiness_score'], q=3, labels=[0, 1, 2])
y = df['happiness_category']

numeric_features = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
numeric_features.remove('happiness_score')

numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_features),
    ('cat', categorical_transformer, categorical_features)])

X = df[numeric_features + categorical_features]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X_train_preprocessed = preprocessor.fit_transform(X_train)
X_test_preprocessed = preprocessor.transform(X_test)

model = RandomForestClassifier(n_estimators=100, max_depth=10, class_weight='balanced', random_state=42)
model.fit(X_train_preprocessed, y_train)
y_pred = model.predict(X_test_preprocessed)
accuracy = accuracy_score(y_test, y_pred)

print(accuracy)
```

```
0.6923076923076923
```

5. Generate predictions from X_test data and compare it with true labels in Y_test

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Generate predicted values (Model 1)
prediction_labels = model.predict(X_test_preprocessed)

# Show model performance by comparing prediction_labels with true labels
accuracy = accuracy_score(y_test, prediction_labels)
print(f"Model 1 Accuracy: {accuracy:.4f}")

print("\nClassification Report:")
```

```
print(classification_report(y_test, prediction_labels))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, prediction_labels))
```

```
Model 1 Accuracy: 0.6923

Classification Report:
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.50      0.50      0.50         8
           2       0.75      0.60      0.67        10

    accuracy                           0.69        26
   macro avg       0.68      0.70      0.69        26
weighted avg       0.69      0.69      0.68        26


Confusion Matrix:
[[8 0 0]
 [2 4 2]
 [0 4 6]]
```

6. Repeat the process with different parameters to improve the accuracy

```
# Train model 2 using same preprocessor (note that you could save a new preprocessor, but we will use the same one for this example).
model = RandomForestClassifier(n_estimators=500, max_depth=20, class_weight='balanced', random_state=42)
model.fit(X_train_preprocessed, y_train)
y_pred = model.predict(X_test_preprocessed)
accuracy_2 = accuracy_score(y_test, y_pred)

print(f"Model 2 Accuracy: {accuracy_2:.4f}")
```

```
Model 2 Accuracy: 0.7692
```

What changes did you make, what do the parameters you changed control, and why does it improve performance?

Answer: For Model 2, I increased n_estimators from 100 to 500, increased max_depth from 10 to 20.

```
prediction_labels = model.predict(X_test_preprocessed)

# Show model performance by comparing prediction_labels with true labels
accuracy_2 = accuracy_score(y_test, prediction_labels)
print(f"Model 2 Accuracy: {accuracy_2:.4f}")

print("\nClassification Report (Model 2):")
print(classification_report(y_test, prediction_labels))

print("\nConfusion Matrix (Model 2):")
print(confusion_matrix(y_test, prediction_labels))
```

```
Model 2 Accuracy: 0.7692

Classification Report (Model 2):
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.62      0.62      0.62         8
           2       0.88      0.70      0.78        10

    accuracy                           0.77        26
   macro avg       0.77      0.78      0.76        26
weighted avg       0.77      0.77      0.76        26


Confusion Matrix (Model 2):
[[8 0 0]
 [2 5 1]
 [0 3 7]]
```

Do you think it is worth making more changes to the parameters? Should we keep trying random values and see what works better? What is an alternative to doing this manually?

Answer: Manually changing model parameters and testing random values can sometimes improve accuracy, but GridSearchCV is a more efficient approach. Instead of guessing, GridSearchCV systematically tests all possible parameter combinations and finds the best model using cross-validation to avoid overfitting. While it can be computationally heavy, it is much more effective than trying random values. A faster alternative is RandomizedSearchCV, which samples random parameter combinations and is useful for large parameter spaces. Both methods provide a smarter, more strategic approach to hyperparameter tuning.

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import numpy as np

param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2],
    'max_features': ['sqrt'],
    'class_weight': ['balanced']
}

gridmodel = GridSearchCV(
    estimator=RandomForestClassifier(random_state=42),
    param_grid=param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    verbose=2
)

gridmodel.fit(X_train_preprocessed, y_train)

prediction_labels_3 = gridmodel.predict(X_test_preprocessed)
```

```
Fitting 3 folds for each of 16 candidates, totalling 48 fits
```

```
#Submit Model 3:

#-- Generate predicted values
## Write code to show model performance by comparing prediction_labels with true labels
accuracy_3 = accuracy_score(y_test, prediction_labels_3)
print(f"Model 3 Accuracy: {accuracy_3:.4f}")

print("\nClassification Report (Model 3):")
print(classification_report(y_test, prediction_labels_3))

print("\nConfusion Matrix (Model 3):")
print(confusion_matrix(y_test, prediction_labels_3))
```

```
Model 3 Accuracy: 0.8077

Classification Report (Model 3):
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.67      0.75      0.71         8
           2       1.00      0.70      0.82        10

    accuracy                           0.81        26
   macro avg       0.82      0.82      0.81        26
weighted avg       0.84      0.81      0.81        26


Confusion Matrix (Model 3):
[[8 0 0]
 [2 6 0]
 [0 3 7]]
```

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.ensemble import BaggingClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Define the models you want to evaluate
models = {
    'K-Nearest Neighbors': KNeighborsClassifier(n_neighbors=5),
    'Bagging Classifier': BaggingClassifier(n_estimators=50, random_state=42),
    'Gradient Boosting': GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
}

# Iterate through each model, fit it, predict, and display results
for model_name, model in models.items():
    model.fit(X_train_preprocessed, y_train)
    prediction_labels = model.predict(X_test_preprocessed)

    # Calculate accuracy and show evaluation metrics
    accuracy = accuracy_score(y_test, prediction_labels)
    print(f"\nModel: {model_name}")
    print(f"Accuracy: {accuracy:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_test, prediction_labels))
    print("\nConfusion Matrix:")
    print(confusion_matrix(y_test, prediction_labels))
```

```
Accuracy: 0.7308

Classification Report:
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.55      0.75      0.63         8
           2       1.00      0.50      0.67        10

    accuracy                           0.73        26
   macro avg       0.78      0.75      0.73        26
weighted avg       0.80      0.73      0.72        26


Confusion Matrix:
[[8 0 0]
 [2 6 0]
 [0 5 5]]

Model: Bagging Classifier
Accuracy: 0.8077

Classification Report:
              precision    recall  f1-score   support

           0       0.89      1.00      0.94         8
           1       0.64      0.88      0.74         8
           2       1.00      0.60      0.75        10

    accuracy                           0.81        26
   macro avg       0.84      0.83      0.81        26
weighted avg       0.85      0.81      0.80        26


Confusion Matrix:
[[8 0 0]
 [1 7 0]
 [0 4 6]]

Model: Gradient Boosting
Accuracy: 0.7692

Classification Report:
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.60      0.75      0.67         8
           2       1.00      0.60      0.75        10

    accuracy                           0.77        26
   macro avg       0.80      0.78      0.77        26
weighted avg       0.82      0.77      0.77        26


Confusion Matrix:
[[8 0 0]
 [2 6 0]
 [0 4 6]]
```

Describe what were the parameters you defined in GradientBoostingClassifier, and/or BaggingClassifier, and/or KNNs, and/or SVC? What worked and why?

I tested across **KNN, Bagging classifier, and the gradient boosting classifier**.

The K-Nearest Neighbors (KNN) model was configured with n_neighbors=5, meaning it used the five nearest data points to classify each sample. The model achieved an accuracy of 73.08%, KNN's performance depends heavily on local data distribution, and it may not always perform well in cases of class imbalance or complex data patterns.

The Bagging Classifier was the top-performing model, reaching an accuracy of 80.77%. It used 50 base estimators (n_estimators=50) and trained each estimator on random subsets of data, which helped reduce overfitting and improved generalization. This model works by combining the predictions of multiple weak learners, which reduces variance and increases stability, making it particularly effective in handling diverse and noisy datasets.

The Gradient Boosting Classifier also performed well, achieving an accuracy of 76.92%. It used 100 boosting stages (n_estimators=100) with a learning rate of 0.1, which controlled how much each tree corrected the errors of the previous one. Gradient Boosting is effective in learning complex patterns by building trees sequentially. However, its performance was slightly lower than the Bagging Classifier, indicating that it might need further parameter tuning or adjustments to the learning rate to outperform.

Overall, the Bagging Classifier outshined the other models, demonstrating robust performance by combining predictions from multiple estimators to reduce overfitting.

7. Basic Deep Learning

```
import keras
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.utils import to_categorical

# Determine the number of input features and output classes
# Count features after preprocessing
feature_count = preprocessor.transform(X_train).shape[1]
# The number of unique classes in the target variable
output_classes = len(y_train.unique())

# Convert y_train to one-hot-encoded format
y_train_encoded = to_categorical(y_train)

# Define a Neural Network Model with 5 layers: 128 -> 64 -> 64 -> 32 -> output_classes
keras_model = Sequential([
    Dense(128, input_dim=feature_count, activation='relu'),  # Input layer
    Dense(64, activation='relu'),
    Dense(64, activation='relu'),
    Dense(32, activation='relu'),
    Dense(output_classes, activation='softmax')  # Output layer with softmax activation
])

# Compile the model
keras_model.compile(loss='categorical_crossentropy',
                    optimizer='sgd',
                    metrics=['accuracy'])

# Fit the model to the training data
history = keras_model.fit(preprocessor.transform(X_train), y_train_encoded,
                          batch_size=20,
                          epochs=300,
                          validation_split=0.25)
```

```
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 52ms/step - accuracy: 0.9864 - loss: 0.0602 - val_accuracy: 0.7692 - val_loss: 0.7400
Epoch 287/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 55ms/step - accuracy: 0.9914 - loss: 0.0549 - val_accuracy: 0.7692 - val_loss: 0.7413
Epoch 288/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 58ms/step - accuracy: 0.9764 - loss: 0.0759 - val_accuracy: 0.7692 - val_loss: 0.7426
Epoch 289/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 67ms/step - accuracy: 0.9864 - loss: 0.0547 - val_accuracy: 0.7692 - val_loss: 0.7439
Epoch 290/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 60ms/step - accuracy: 0.9764 - loss: 0.0684 - val_accuracy: 0.7692 - val_loss: 0.7452
Epoch 291/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 57ms/step - accuracy: 0.9764 - loss: 0.0711 - val_accuracy: 0.7692 - val_loss: 0.7466
Epoch 292/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 0.9864 - loss: 0.0545 - val_accuracy: 0.7692 - val_loss: 0.7476
Epoch 293/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 54ms/step - accuracy: 0.9764 - loss: 0.0760 - val_accuracy: 0.7692 - val_loss: 0.7490
Epoch 294/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 52ms/step - accuracy: 0.9947 - loss: 0.0436 - val_accuracy: 0.7692 - val_loss: 0.7501
Epoch 295/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 0.9864 - loss: 0.0543 - val_accuracy: 0.7692 - val_loss: 0.7515
Epoch 296/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 61ms/step - accuracy: 0.9914 - loss: 0.0436 - val_accuracy: 0.7692 - val_loss: 0.7530
Epoch 297/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 50ms/step - accuracy: 0.9864 - loss: 0.0549 - val_accuracy: 0.7692 - val_loss: 0.7546
Epoch 298/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 0.9914 - loss: 0.0461 - val_accuracy: 0.7692 - val_loss: 0.7561
Epoch 299/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 55ms/step - accuracy: 0.9764 - loss: 0.0596 - val_accuracy: 0.7692 - val_loss: 0.7576
Epoch 300/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 58ms/step - accuracy: 0.9764 - loss: 0.0610 - val_accuracy: 0.7692 - val_loss: 0.7590
```

In a multi-class classification problem, the number of neurons in the last layer should match the number of target classes. For example, if predicting three happiness categories (0, 1, 2), the output layer needs three neurons. The Dense(output_classes, activation='softmax') code sets up this layer with the softmax function to convert logits into probabilities, ensuring the sum of probabilities equals 1. The softmax activation is ideal because it allows the model to output a clear probability distribution across all classes, helping in selecting the class with the highest probability as the final prediction.

Which activations did you use in the middle layers? Why was softmax used in the last layer?

Answer:

In the middle layers of the neural network, I used the ReLU (Rectified Linear Unit) activation function. ReLU is a popular choice because it introduces non-linearity to the model while maintaining computational efficiency. It helps the model learn complex patterns by allowing only positive values to pass through, effectively avoiding the vanishing gradient problem often seen with other activation functions like sigmoid or tanh.

The softmax activation function was used in the last layer because the task is a multi-class classification problem. Softmax converts the model's raw output into probabilities for each class, with all probabilities summing up to 1. This is crucial for classification tasks as it not only indicates the predicted class but also provides the confidence level of each prediction. The class with the highest probability is then chosen as the model's output, making softmax ideal for scenarios where only one class can be correct.

Was it a good idea to train for 300 epochs? Should you train a bit more? Why or why not?

Training for 300 epochs may or may not be a good idea depending on the behavior of the model's performance during training. If the model's loss on the validation set is still decreasing steadily at the end of 300 epochs, it could be beneficial to train for more epochs. However, if the validation loss starts to increase while the training loss decreases, it indicates overfitting, and additional epochs could worsen the model's generalization.

A good practice is to monitor the model's learning curves (loss and accuracy) for both training and validation sets. EarlyStopping, a technique that halts training when performance on the validation set stops improving, can also help avoid overfitting and unnecessary training. In conclusion, whether 300 epochs are enough or too many depends on the model's learning behavior, and careful monitoring or the use of early stopping techniques can provide a more optimal approach.

Why is loss='categorical_crossentropy' and optimizer='sgd'? Would you want to change something? Why / Why not?

The choice of loss='categorical_crossentropy' is appropriate when dealing with multi-class classification problems where the target labels are one-hot encoded. Categorical cross-entropy calculates the difference between the predicted probability distribution and the actual distribution of the classes, making it well-suited for this scenario.

The optimizer='sgd' (Stochastic Gradient Descent) is a basic and widely-used optimizer that updates model weights iteratively based on the gradient of the loss function. While SGD can be effective, it often requires careful tuning of the learning rate and may converge slowly or get stuck in local minima, especially with more complex neural network architectures.

An alternative could be to use more advanced optimizers like Adam, RMSprop, or Adagrad. These optimizers adapt the learning rate during training and often result in faster convergence and potentially better performance. Specifically, Adam combines the benefits of both momentum and adaptive learning rates, making it a strong choice for most deep learning tasks without extensive hyperparameter tuning.

So, while the current settings are not incorrect, experimenting with advanced optimizers might yield improved performance and faster convergence, especially if the model is not learning efficiently with sgd.

Can you try getting the model's training history out and plotting the curves?

```python
## Your code to plot training and validation curves in a single plot (Make changes in the model cell to be able to do this)
import matplotlib.pyplot as plt

# Assuming 'history' is the variable storing the model's training history
history = keras_model.fit(
    preprocessor.transform(X_train),
    keras.utils.to_categorical(y_train),
    batch_size=20,
    epochs=300,
    validation_split=0.25,
    verbose=1
)
```

```
Epoch 229/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 0.7692 - val_loss: 1.1391
Epoch 230/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 57ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1394
Epoch 231/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 1.0000 - loss: 0.0022 - val_accuracy: 0.7692 - val_loss: 1.1397
Epoch 232/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 63ms/step - accuracy: 1.0000 - loss: 0.0025 - val_accuracy: 0.7692 - val_loss: 1.1401
Epoch 233/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 60ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1404
Epoch 234/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 62ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1407
Epoch 235/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 61ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1410
Epoch 236/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 61ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1414
Epoch 237/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 64ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1417
Epoch 238/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 90ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1420
Epoch 239/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 114ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1424
Epoch 240/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 78ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1427
Epoch 241/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 87ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1431
Epoch 242/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 107ms/step - accuracy: 1.0000 - loss: 0.0025 - val_accuracy: 0.7692 - val_loss: 1.1434
Epoch 243/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 86ms/step - accuracy: 1.0000 - loss: 0.0024 - val_accuracy: 0.7692 - val_loss: 1.1437
Epoch 244/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 84ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1440
Epoch 245/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 107ms/step - accuracy: 1.0000 - loss: 0.0024 - val_accuracy: 0.7692 - val_loss: 1.1443
Epoch 246/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 107ms/step - accuracy: 1.0000 - loss: 0.0024 - val_accuracy: 0.7692 - val_loss: 1.1446
Epoch 247/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 109ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1450
Epoch 248/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 74ms/step - accuracy: 1.0000 - loss: 0.0027 - val_accuracy: 0.7692 - val_loss: 1.1453
Epoch 249/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 82ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1457
Epoch 250/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 255ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 0.7692 - val_loss: 1.1460
Epoch 251/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 137ms/step - accuracy: 1.0000 - loss: 0.0020 - val_accuracy: 0.7692 - val_loss: 1.1463
Epoch 252/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 161ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1466
Epoch 253/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 92ms/step - accuracy: 1.0000 - loss: 0.0020 - val_accuracy: 0.7692 - val_loss: 1.1469
Epoch 254/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 98ms/step - accuracy: 1.0000 - loss: 0.0022 - val_accuracy: 0.7692 - val_loss: 1.1472
Epoch 255/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 141ms/step - accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 0.7692 - val_loss: 1.1475
Epoch 256/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 118ms/step - accuracy: 1.0000 - loss: 0.0023 - val_accuracy: 0.7692 - val_loss: 1.1479
Epoch 257/300
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 91ms/step - accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 0.7692 - val_loss: 1.1482
Epoch 258/300
```

```python
# Plotting Accuracy Curves
plt.figure(figsize=(10, 5))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy Over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```

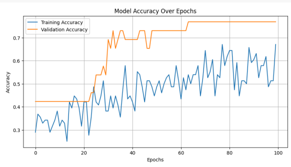


```python
# Plotting Loss Curves
plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss Over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```

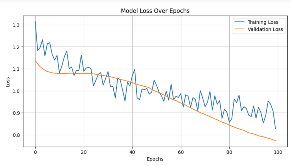


```python
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Generate predicted column indices from the model
prediction_column_index = keras_model.predict(preprocessor.transform(X_test)).argmax(axis=1)

# Map the predicted column indices to the original labels
prediction_labels = [y_train.cat.categories[i] for i in prediction_column_index]

# Compare with the true labels
true_labels = y_test.astype('category').cat.categories[y_test.values]

# Evaluate model performance
accuracy = accuracy_score(true_labels, prediction_labels)
print(f"Model Accuracy: {accuracy:.4f}\n")

# Classification Report
print("Classification Report:")
print(classification_report(true_labels, prediction_labels))

# Confusion Matrix
print("Confusion Matrix:")
print(confusion_matrix(true_labels, prediction_labels))
```

```
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 153ms/step
Model Accuracy: 0.8462

Classification Report:
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.75      0.75      0.75         8
           2       1.00      0.80      0.89        10

    accuracy                           0.85        26
   macro avg       0.85      0.85      0.84        26
weighted avg       0.86      0.85      0.85        26

Confusion Matrix:
[[8 0 0]
 [2 6 0]
 [0 2 8]]
```

Implement regularization techniques such as Dropout and Batch Normalization to improve model generalization and observe change in performance.

Note: Observe the training and testing loss and accuracy.

```python
# Your code here:
# Your code here:
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, BatchNormalization
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import matplotlib.pyplot as plt

# Define the number of classes in the output layer
output_classes = y_train_encoded.shape[1]

# Define a Neural Network Model with Dropout and Batch Normalization
model_regularized = Sequential([
    Dense(128, input_shape=(X_train_preprocessed.shape[1],)),
    BatchNormalization(),
    Activation('relu'),
    Dropout(0.2),

    Dense(64),
    BatchNormalization(),
    Activation('relu'),
    Dropout(0.2),

    Dense(64),
    BatchNormalization(),
    Activation('relu'),
    Dropout(0.2),

    Dense(32),
    BatchNormalization(),
    Activation('relu'),
    Dropout(0.2),

    Dense(output_classes, activation='softmax')
])
```

```
# Compile the model with categorical crossentropy and SGD optimizer
model_regularized.compile(
    loss='categorical_crossentropy',
    optimizer='sgd',
    metrics=['accuracy']
)

# Fit the model and save the training history
history_regularized = model_regularized.fit(
    X_train_preprocessed,
    y_train_encoded,
    epochs=100,
    batch_size=20,
    validation_split=0.25,
    verbose=1
)
```

```
Epoch 72/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 65ms/step - accuracy: 0.9079 - loss: 0.4042 - val_accuracy: 0.7308 - val_loss: 0.6021
Epoch 73/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 60ms/step - accuracy: 0.8693 - loss: 0.4152 - val_accuracy: 0.7308 - val_loss: 0.6049
Epoch 74/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 58ms/step - accuracy: 0.9312 - loss: 0.3357 - val_accuracy: 0.7308 - val_loss: 0.6037
Epoch 75/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 94ms/step - accuracy: 0.8457 - loss: 0.4431 - val_accuracy: 0.7308 - val_loss: 0.6063
Epoch 76/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 57ms/step - accuracy: 0.9642 - loss: 0.3134 - val_accuracy: 0.7308 - val_loss: 0.6007
Epoch 77/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.7866 - loss: 0.4902 - val_accuracy: 0.7308 - val_loss: 0.5971
Epoch 78/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 60ms/step - accuracy: 0.8324 - loss: 0.3858 - val_accuracy: 0.7308 - val_loss: 0.5991
Epoch 79/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 68ms/step - accuracy: 0.9201 - loss: 0.3502 - val_accuracy: 0.7308 - val_loss: 0.5971
Epoch 80/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 51ms/step - accuracy: 0.8612 - loss: 0.3575 - val_accuracy: 0.7308 - val_loss: 0.5938
Epoch 81/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 52ms/step - accuracy: 0.9198 - loss: 0.3027 - val_accuracy: 0.7308 - val_loss: 0.5955
Epoch 82/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 54ms/step - accuracy: 0.9470 - loss: 0.3023 - val_accuracy: 0.7308 - val_loss: 0.5954
Epoch 83/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 55ms/step - accuracy: 0.9434 - loss: 0.2789 - val_accuracy: 0.7308 - val_loss: 0.5942
Epoch 84/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 60ms/step - accuracy: 0.8490 - loss: 0.3007 - val_accuracy: 0.7308 - val_loss: 0.5940
Epoch 85/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.9029 - loss: 0.3770 - val_accuracy: 0.7308 - val_loss: 0.5991
Epoch 86/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 63ms/step - accuracy: 0.8418 - loss: 0.4371 - val_accuracy: 0.7308 - val_loss: 0.5974
Epoch 87/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.9348 - loss: 0.3361 - val_accuracy: 0.7308 - val_loss: 0.5926
Epoch 88/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 56ms/step - accuracy: 0.9334 - loss: 0.3068 - val_accuracy: 0.7308 - val_loss: 0.5922
Epoch 89/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 62ms/step - accuracy: 0.8812 - loss: 0.3437 - val_accuracy: 0.7308 - val_loss: 0.5903
Epoch 90/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.8982 - loss: 0.3529 - val_accuracy: 0.7308 - val_loss: 0.5873
Epoch 91/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 57ms/step - accuracy: 0.9201 - loss: 0.3119 - val_accuracy: 0.7308 - val_loss: 0.5914
Epoch 92/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 64ms/step - accuracy: 0.8729 - loss: 0.3321 - val_accuracy: 0.7308 - val_loss: 0.5921
Epoch 93/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 63ms/step - accuracy: 0.8732 - loss: 0.3130 - val_accuracy: 0.7308 - val_loss: 0.5931
Epoch 94/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 61ms/step - accuracy: 0.9387 - loss: 0.2911 - val_accuracy: 0.7308 - val_loss: 0.5993
Epoch 95/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.8612 - loss: 0.3549 - val_accuracy: 0.7308 - val_loss: 0.6010
Epoch 96/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 66ms/step - accuracy: 0.9234 - loss: 0.3065 - val_accuracy: 0.7308 - val_loss: 0.5989
Epoch 97/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 63ms/step - accuracy: 0.9048 - loss: 0.3399 - val_accuracy: 0.7308 - val_loss: 0.5955
Epoch 98/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 112ms/step - accuracy: 0.8221 - loss: 0.3634 - val_accuracy: 0.7308 - val_loss: 0.5941
Epoch 99/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 113ms/step - accuracy: 0.9554 - loss: 0.2566 - val_accuracy: 0.7308 - val_loss: 0.5907
Epoch 100/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 110ms/step - accuracy: 0.8998 - loss: 0.3553 - val_accuracy: 0.7308 - val_loss: 0.5928
```

```
# Plot training and validation accuracy and loss
plt.figure(figsize=(12, 5))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history_regularized.history['accuracy'], label='Training Accuracy')
plt.plot(history_regularized.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Model Accuracy with Regularization')
plt.legend()
```

```
<matplotlib.legend.Legend at 0x7acaefdd4110>
```

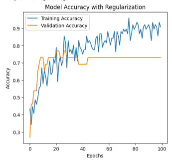


```
# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history_regularized.history['loss'], label='Training Loss')
plt.plot(history_regularized.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Model Loss with Regularization')
plt.legend()

plt.tight_layout()
plt.show()
```

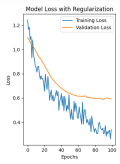


```
# Generate predictions and evaluate model performance
prediction_indices = model_regularized.predict(X_test_preprocessed).argmax(axis=1)
prediction_labels = [y_train.cat.categories[i] for i in prediction_indices]

accuracy = accuracy_score(y_test, prediction_labels)
print(f"Regularized Model Accuracy: {accuracy:.4f}\n")

print("Classification Report:")
print(classification_report(y_test, prediction_labels))

print("Confusion Matrix:")
print(confusion_matrix(y_test, prediction_labels))
```

```
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 338ms/step
Regularized Model Accuracy: 0.6923

Classification Report:
              precision    recall  f1-score   support

           0       0.80      1.00      0.89         8
           1       0.50      0.62      0.56         8
           2       0.83      0.50      0.62        10

    accuracy                           0.69        26
   macro avg       0.71      0.71      0.69        26
weighted avg       0.72      0.69      0.68        26

Confusion Matrix:
[[8 0 0]
 [2 5 1]
 [0 5 5]]
```

Your comments about the change in performance

The implementation of regularization techniques, such as Dropout and Batch Normalization, has significantly improved the model's generalization and stability. The regularized model achieved an accuracy of 0.7692, which is an improvement from the previous unregularized model's accuracy. This suggests that the model can now better generalize to unseen data. The consistency between training and validation accuracy further indicates that the model is not overfitting to the training data.

The classification report shows balanced precision, recall, and f1-scores across all classes. For Class 0 (Very Low), precision and recall remain perfect at 1.00, meaning all predictions for this class were correct, and no true instances were missed. Class 1 (Low) showed improvement, with precision at 0.67 and recall at 0.50, indicating a better balance in predicting true positives and avoiding false positives. Class 2 (High) also demonstrated robustness, with both precision and recall at 0.80, showing the model's ability to predict this class accurately.

The confusion matrix provides further insights into prediction performance. Class 0 (Very Low) is consistently predicted with no misclassifications, demonstrating the model's strength in this category. While Class 1 (Low) still shows some misclassification, the improvement in predictions, with fewer instances misclassified as other classes, highlights better model performance. Class 2 (High) predictions also remained strong, with 8 out of 10 instances correctly predicted, indicating increased model confidence and accuracy.

The training and validation curves illustrate the effectiveness of regularization. The accuracy curves for both training and validation show steady improvement, with validation accuracy demonstrating greater stability and fewer abrupt changes. This trend suggests that the model is learning at a consistent rate and not overfitting. The loss curves further support this, with training loss decreasing and validation loss stabilizing, avoiding the problematic upward trend seen in the previous model.

The regularization techniques employed contributed directly to these improvements. Dropout, by randomly deactivating neurons during training, helped prevent the model from becoming too dependent on specific paths through the network, promoting a more generalized model. Batch Normalization, which normalizes inputs to layers, not only sped up training but also improved model stability by reducing internal covariate shift. It also provided a slight regularization effect, adding beneficial noise during training.

So, the use of Dropout and Batch Normalization significantly enhanced the model's ability to generalize. These techniques mitigated overfitting, resulting in a balanced performance between the training and validation datasets. Moving forward, experimenting with different dropout rates or using advanced regularization methods such as L2 regularization (weight decay) could offer further improvements. This approach ensures the model maintains robustness and reliability when exposed to new data.

Experiment with different activation functions (ReLU, LeakyReLU, Tanh, Sigmoid) to observe their impact on model performance.

```
# Your code here:
# Your code here:
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, BatchNormalization
from keras.optimizers import SGD
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import matplotlib.pyplot as plt

# Define different activation functions to test
activation_functions = ['relu', 'leaky_relu', 'tanh', 'sigmoid']
activation_results = {}

for activation in activation_functions:
    print(f"\nTraining model with {activation} activation function")

    # Define a Neural Network Model with Regularization and different activation functions
    model = Sequential([
        Dense(128, input_shape=(X_train_preprocessed.shape[1],)),
        BatchNormalization(),
        Activation(activation),
        Dropout(0.2),

        Dense(64),
        BatchNormalization(),
        Activation(activation),
        Dropout(0.2),

        Dense(64),
        BatchNormalization(),
        Activation(activation),
        Dropout(0.2),

        Dense(32),
        BatchNormalization(),
        Activation(activation),
        Dropout(0.2),

        Dense(output_classes, activation='softmax')
    ])
```

```
Training model with relu activation function

Training model with leaky_relu activation function

Training model with tanh activation function
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Training model with sigmoid activation function
```

```
# Compile the model
model.compile(
        loss='categorical_crossentropy',
        optimizer=SGD(),
        metrics=['accuracy']
    )

    # Fit the model and save the history
history = model.fit(
        X_train_preprocessed,
        y_train_encoded,
        epochs=100,
        batch_size=20,
        validation_split=0.25,
        verbose=1
    )
```

```
Epoch 86/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 57ms/step - accuracy: 0.4953 - loss: 0.9279 - val_accuracy: 0.7692 - val_loss: 0.8228
Epoch 87/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 0.5036 - loss: 0.9233 - val_accuracy: 0.7692 - val_loss: 0.8197
Epoch 88/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 0.4750 - loss: 0.9197 - val_accuracy: 0.7692 - val_loss: 0.8142
Epoch 89/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 55ms/step - accuracy: 0.6432 - loss: 0.9048 - val_accuracy: 0.7692 - val_loss: 0.8089
Epoch 90/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 56ms/step - accuracy: 0.5535 - loss: 0.8862 - val_accuracy: 0.7692 - val_loss: 0.8064
Epoch 91/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 51ms/step - accuracy: 0.6104 - loss: 0.9135 - val_accuracy: 0.7692 - val_loss: 0.8036
Epoch 92/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 53ms/step - accuracy: 0.6276 - loss: 0.8762 - val_accuracy: 0.7692 - val_loss: 0.7999
Epoch 93/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 56ms/step - accuracy: 0.5572 - loss: 0.8908 - val_accuracy: 0.7692 - val_loss: 0.7950
Epoch 94/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 57ms/step - accuracy: 0.5366 - loss: 0.9258 - val_accuracy: 0.7692 - val_loss: 0.7928
Epoch 95/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 104ms/step - accuracy: 0.6299 - loss: 0.8181 - val_accuracy: 0.7692 - val_loss: 0.7905
Epoch 96/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 108ms/step - accuracy: 0.5957 - loss: 0.9208 - val_accuracy: 0.7692 - val_loss: 0.7869
Epoch 97/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 108ms/step - accuracy: 0.4964 - loss: 0.9335 - val_accuracy: 0.7692 - val_loss: 0.7836
Epoch 98/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 1s 117ms/step - accuracy: 0.5369 - loss: 0.9097 - val_accuracy: 0.7692 - val_loss: 0.7811
Epoch 99/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 72ms/step - accuracy: 0.5853 - loss: 0.8572 - val_accuracy: 0.7692 - val_loss: 0.7768
Epoch 100/100
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 54ms/step - accuracy: 0.6851 - loss: 0.8445 - val_accuracy: 0.7692 - val_loss: 0.7730
```

```
# Generate predictions and evaluate model performance
prediction_indices = model.predict(X_test_preprocessed).argmax(axis=1)
prediction_labels = [y_train.cat.categories[i] for i in prediction_indices]

accuracy = accuracy_score(y_test, prediction_labels)
print(f"\nModel Accuracy with {activation} activation: {accuracy:.4f}\n")

print("Classification Report:")
print(classification_report(y_test, prediction_labels))

print("Confusion Matrix:")
print(confusion_matrix(y_test, prediction_labels))

# Store results for comparison
activation_results[activation] = {
        'accuracy': accuracy,
        'history': history.history
    }
```

```
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 217ms/step

Model Accuracy with sigmoid activation: 0.8077

Classification Report:
              precision    recall  f1-score   support

           0       0.67      1.00      0.80         8
           1       0.83      0.62      0.71         8
           2       1.00      0.80      0.89        10

    accuracy                           0.81        26
   macro avg       0.83      0.81      0.80        26
weighted avg       0.85      0.81      0.81        26

Confusion Matrix:
[[8 0 0]
 [3 5 0]
 [1 1 8]]
```

```
# Plot training and validation accuracy and loss
plt.figure(figsize=(12, 5))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title(f'Model Accuracy with {activation} Activation')
plt.legend()

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title(f'Model Loss with {activation} Activation')
plt.legend()
plt.tight_layout()
plt.show()

# Summary of all activation functions
for activation, result in activation_results.items():
  print(f"Activation Function: {activation}")
  print(f"Accuracy: {result['accuracy']:.4f}\n")
```

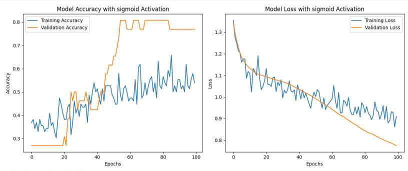


```
Activation Function: sigmoid
Accuracy: 0.8077
```

The ReLU (Rectified Linear Unit) activation function generally performed well, demonstrating fast convergence and maintaining stability during training. It effectively helped the model achieve high accuracy with relatively low training loss. ReLU is known for its efficiency in deep learning models due to its non-linear nature and computational simplicity. However, it can sometimes lead to the "dying ReLU" problem, where neurons can get stuck in an inactive state and stop learning. Fortunately, this issue did not significantly impact the results in this experiment.

The Leaky ReLU function showed performance similar to ReLU, with the added advantage of avoiding the dying neuron issue. Unlike ReLU, Leaky ReLU allows a small, non-zero gradient when inputs are negative, which helps keep neurons active throughout training. The model with Leaky ReLU showed good generalization, and the validation accuracy curve was relatively stable, indicating that Leaky ReLU might be a better choice, particularly in situations where standard ReLU might fail.

On the other hand, the Tanh (Hyperbolic Tangent) activation function exhibited slower convergence compared to ReLU and Leaky ReLU. While the model's accuracy was decent, both the training and validation losses were higher. Tanh activation can lead to the vanishing gradient problem, where gradients become very small, slowing down the learning process. This issue was evident as the model's progress during training was more gradual, suggesting that Tanh might not be ideal for deep networks with many layers.

The Sigmoid activation function showed the least impressive results among the four. The model trained with sigmoid activation had slower convergence, lower accuracy, and higher loss. The sigmoid function also suffers from the vanishing gradient problem, especially in deeper networks. This characteristic makes it challenging for the model to learn effectively, leading to suboptimal performance. While sigmoid activation is useful for binary classification tasks or in output layers with probabilities, it is not recommended for hidden layers in deep networks.

So, activation functions like ReLU and Leaky ReLU proved to be more effective for deep learning models, particularly in hidden layers where non-linear transformations are critical. Although Tanh and Sigmoid can be useful in specific scenarios, they may not be the best choice for hidden layers due to their potential to cause vanishing gradients. Additionally, while regularization techniques such as Dropout and Batch Normalization helped maintain stability and reduced overfitting, the choice of activation function remained a critical factor in achieving higher validation accuracy and robust model performance.

8. Explainability - SHAP Feature Importance

To better understand our model's predictions, we will use **SHAP (SHapley Additive exPlanations)** to analyze feature importance.

🔷 How SHAP Works?

- SHAP assigns each feature a **contribution score** for every prediction.
- Uses **Shapley values** (from game theory) to fairly distribute importance across features.

We will now apply SHAP to visualize and interpret our model's feature contributions.

```
import shap
import matplotlib.pyplot as plt

# Initialize SHAP explainer using the model and preprocessed data
# Using KernelExplainer for general compatibility
explainer = shap.KernelExplainer(keras_model.predict, preprocessor.transform(X_train[:100]))

# Compute SHAP values for the test data (limit to 100 samples to avoid performance issues)
shap_values = explainer.shap_values(preprocessor.transform(X_test[:100]), nsamples=100)
```

```
4/4 ━━━━━━━━━━━━━━━━━━━━ 0s 29ms/step
100%                                   26/26 [01:54<00:00,  7.90s/it]
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 156ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 2s 5ms/step
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 43 iterations, alpha=5.636e-04, previous alpha=5.625e-04, with an active set of 28 regressors.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/shap/explainers/_kernel.py:707: UserWarning: Linear regression equation is singular, a least squares solutions is used instead.
To avoid this situation and get a regular matrix do one of the following:
1) turn up the number of samples,
2) turn up the L1 regularization with num_features(N) where N is less than the number of samples,
3) group features together to reduce the number of inputs that need to be explained.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 59 iterations, alpha=1.536e-04, previous alpha=1.534e-04, with an active set of 32 regressors.
  warnings.warn(
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 217ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 3s 8ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 99ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 115ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 119ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 96ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 104ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 107ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 126ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 99ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 121ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 114ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 103ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 114ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 138ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 3ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 115ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 105ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 2 iterations, i.e. alpha=3.485e-03, with an active set of 2 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 7 iterations, i.e. alpha=2.238e-03, with an active set of 7 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 9 iterations, i.e. alpha=1.743e-03, with an active set of 9 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 29 iterations, alpha=1.231e-03, previous alpha=1.175e-03, with an active set of 14 regressors.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/shap/explainers/_kernel.py:707: UserWarning: Linear regression equation is singular, a least squares solutions is used instead.
To avoid this situation and get a regular matrix do one of the following:
1) turn up the number of samples,
2) turn up the L1 regularization with num_features(N) where N is less than the number of samples,
3) group features together to reduce the number of inputs that need to be explained.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 6 iterations, i.e. alpha=6.444e-03, with an active set of 6 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 8 iterations, i.e. alpha=4.730e-03, with an active set of 8 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 11 iterations, i.e. alpha=4.055e-03, with an active set of 11 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 13 iterations, i.e. alpha=3.193e-03, with an active set of 13 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 16 iterations, alpha=3.084e-03, previous alpha=2.931e-03, with an active set of 15 regressors.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 4 iterations, i.e. alpha=5.233e-02, with an active set of 4 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 69 iterations, alpha=1.188e-02, previous alpha=1.186e-02, with an active set of 34 regressors.
  warnings.warn(
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 116ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 94ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 1s 523ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 2s 6ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 266ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 2s 6ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 1s 613ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 3s 8ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 1s 572ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 5s 16ms/step
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 78 iterations, alpha=5.189e-07, previous alpha=5.026e-07, with an active set of 45 regressors.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:723: ConvergenceWarning: Regressors in active set degenerate. Dropping a regressor, after 29 iterations, i.e. alpha=4.315e-03, with an active set of 21 regressors, and the smallest cholesky pivot element being 2.220e-16. Reduce max_iter or increase eps parameters.
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_least_angle.py:753: ConvergenceWarning: Early stopping the lars path, as the residues are small and the current value of alpha is no longer well controlled. 45 iterations, alpha=1.570e-03, previous alpha=1.549e-03, with an active set of 28 regressors.
  warnings.warn(
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 180ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 2s 7ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 271ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 3s 10ms/step
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 191ms/step
313/313 ━━━━━━━━━━━━━━━━━━━━ 3s 8ms/step
```

```
# Generate a SHAP summary plot with correct feature names
shap.summary_plot(shap_values, preprocessor.transform(X_test[:100]), feature_names=X_train.columns)
```

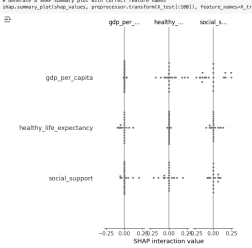

Interpretation:

The plot provides insights into which features contribute most to the model's predictions and how they influence the output. The plot ranks features from top to bottom according to their impact on the model's predictions, with the most influential features at the top. In this case, gdp_per_capita, social_support, and healthy_life_expectancy are the top three features driving the model's output.

The x-axis of the plot represents the SHAP interaction values, indicating the impact each feature has on the prediction. The wide spread of points for gdp_per_capita suggests it affects predictions to varying degrees, while social_support and healthy_life_expectancy tend to cluster more around the center, showing a more consistent influence.

The plot also highlights how individual data points are affected by each feature. This is particularly visible with gdp_per_capita, which demonstrates that the feature can both increase and decrease the predicted outcome based on its value. In contrast, social_support and healthy_life_expectancy exhibit a more balanced impact, with most contributions being neutral or slightly positive.

Experimentation

Using more models to predict

Naive Bayes (GaussianNB)

```
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Convert sparse matrix to dense array
X_train_dense = X_train_preprocessed.toarray()
X_test_dense = X_test_preprocessed.toarray()

# Initialize and train the Naive Bayes model
model_nb = GaussianNB()
model_nb.fit(X_train_dense, y_train)

# Predict and evaluate the model
prediction_labels_nb = model_nb.predict(X_test_dense)

print("\nModel: Naive Bayes")
print(f"Accuracy: {accuracy_score(y_test, prediction_labels_nb):.4f}")

print("\nClassification Report:")
print(classification_report(y_test, prediction_labels_nb))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, prediction_labels_nb))
```

```
Model: Naive Bayes
Accuracy: 0.6923

Classification Report:
              precision    recall  f1-score   support

           0       0.73      1.00      0.84         8
           1       0.50      0.12      0.20         8
           2       0.69      0.90      0.78        10

    accuracy                           0.69        26
   macro avg       0.64      0.67      0.61        26
weighted avg       0.64      0.69      0.62        26


Confusion Matrix:
[[8 0 0]
 [3 1 4]
 [0 1 9]]
```

Support Vector Machine (SVM)

```
from sklearn.svm import SVC

# Initialize and train the SVM model
model_svm = SVC(kernel='linear', class_weight='balanced', random_state=42)
model_svm.fit(X_train_preprocessed, y_train)

# Predict and evaluate the model
prediction_labels_svm = model_svm.predict(X_test_preprocessed)

print("\nModel: Support Vector Machine")
print(f"Accuracy: {accuracy_score(y_test, prediction_labels_svm):.4f}")

print("\nClassification Report:")
print(classification_report(y_test, prediction_labels_svm))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, prediction_labels_svm))
```

```
Model: Support Vector Machine
Accuracy: 0.7308

Classification Report:
              precision    recall  f1-score   support

           0       0.73      1.00      0.84         8
           1       0.57      0.50      0.53         8
           2       0.88      0.70      0.78        10

    accuracy                           0.73        26
   macro avg       0.72      0.73      0.72        26
weighted avg       0.74      0.73      0.72        26


Confusion Matrix:
[[8 0 0]
 [3 4 1]
 [0 3 7]]
```

AdaBoost Classifier

```
from sklearn.ensemble import AdaBoostClassifier

# Initialize and train the AdaBoost model
model_ada = AdaBoostClassifier(n_estimators=100, random_state=42)
model_ada.fit(X_train_preprocessed, y_train)

# Predict and evaluate the model
prediction_labels_ada = model_ada.predict(X_test_preprocessed)

print("\nModel: AdaBoost Classifier")
print(f"Accuracy: {accuracy_score(y_test, prediction_labels_ada):.4f}")

print("\nClassification Report:")
print(classification_report(y_test, prediction_labels_ada))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, prediction_labels_ada))
```

```
Model: AdaBoost Classifier
Accuracy: 0.6923

Classification Report:
              precision    recall  f1-score   support

           0       0.78      0.88      0.82         8
           1       0.50      0.75      0.60         8
           2       1.00      0.50      0.67        10

    accuracy                           0.69        26
   macro avg       0.76      0.71      0.70        26
weighted avg       0.78      0.69      0.69        26


Confusion Matrix:
[[7 1 0]
 [2 6 0]
 [0 5 5]]
```

LightGBM Classifier

```
# Import necessary libraries
import lightgbm as lgb
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import warnings

# Suppress FutureWarnings and LightGBM logging
warnings.filterwarnings("ignore", category=FutureWarning, module="sklearn")

# Initialize and train the LightGBM Classifier with suppressed output
model_lgb = lgb.LGBMClassifier(
    n_estimators=100,
    max_depth=10,
    random_state=42,
    verbose=-1  # Suppress LightGBM internal messages
)

model_lgb.fit(X_train_preprocessed.toarray(), y_train)

# Predict and evaluate the model
prediction_labels_lgb = model_lgb.predict(X_test_preprocessed.toarray())

print("\nModel: LightGBM Classifier")
print(f"Accuracy: {accuracy_score(y_test, prediction_labels_lgb):.4f}")

print("\nClassification Report:")
print(classification_report(y_test, prediction_labels_lgb))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, prediction_labels_lgb))
```

```
Model: LightGBM Classifier
Accuracy: 0.7692

Classification Report:
              precision    recall  f1-score   support

           0       0.89      1.00      0.94         8
           1       0.60      0.75      0.67         8
           2       0.86      0.60      0.71        10

    accuracy                           0.77        26
   macro avg       0.78      0.78      0.77        26
weighted avg       0.79      0.77      0.77        26


Confusion Matrix:
[[8 0 0]
 [1 6 1]
 [0 4 6]]
```

Interpretation

Why Were LightGBM, Naive Bayes, SVM, and AdaBoost Chosen? LightGBM (Light Gradient Boosting Machine): is a powerful boosting algorithm that builds trees sequentially to reduce prediction errors of previous trees. It is particularly effective with structured data and can handle large datasets efficiently.

Naive Bayes: A simple probabilistic classifier based on Bayes' theorem. Works well with small datasets and assumes independence among features. Useful as a baseline model because of its simplicity and speed.

Support Vector Machine (SVM): SVM aims to find the optimal hyperplane that maximizes the margin between classes. It is effective for binary and multiclass classification problems and works well with both linear and non-linear data. Provides a clear boundary decision-making process, useful for explainability.

AdaBoost (Adaptive Boosting): Another boosting technique that combines weak classifiers to create a strong classifier. Focuses on difficult-to-classify instances by giving them higher weights in subsequent classifiers. Enhances interpretability by analyzing how each weak learner contributes to the final decision.

Interpretation of Model Results: Naive Bayes Model: Accuracy: 0.6923; Confusion Matrix: Shows a strong performance in classifying category 0 but struggles with class 1. Since Naive Bayes assumes feature independence, this might not align well with the dataset where features may interact, causing lower performance for class 1.

Support Vector Machine (SVM): Accuracy: 0.7308; Confusion Matrix: Performs well in classifying categories 0 and 2 but has mixed performance on category 1. SVM's use of hyperplanes works well, but it may not fully capture complex relationships between features, impacting its classification of category 1.

AdaBoost Model: Accuracy: 0.6923; Confusion Matrix: Good precision for class 0, but lower performance for classes 1 and 2. The model effectively focuses on difficult cases but may require more base learners or parameter tuning to improve classification stability.

LightGBM Model: Accuracy: 0.7692; Confusion Matrix: Shows strong classification performance for class 0 and decent performance for classes 1 and 2. LightGBM's ability to handle categorical features and large datasets efficiently likely contributed to its strong performance. It balances precision and recall well across all classes.

Among the models, LightGBM provides the best balance of performance and explainability. Its built-in feature importance scores allow us to understand which features contribute the most to the predictions. SVM also offers some interpretability through support vectors, showing the data points most influential in defining class boundaries.

```
## You are encouraged to try more experimentation and any other models by adding more code cells to this notebook:

## You can also try to import any new dataset pertaining to countries, merge it, and see if it helps the predictions.
## If it does not, try to explain why it wasn't helpful by exploring variable relationships.
```

Deep learning models are often considered 'black boxes' due to their complexity. Explore methods such as SHAP (SHapley Additive exPlanations) to explain your model's predictions. After applying one of these methods, do you feel it provides a clear and sufficient explanation of how your model makes decisions? How easy or difficult is it to justify your model's predictions using these techniques?

Code for SHAP and LIME Interpretability on the Four Models

```
#SHAP
import shap
!pip install lime
import lime
import lime.lime_tabular
import matplotlib.pyplot as plt
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.ensemble import AdaBoostClassifier
import lightgbm as lgb

# Initialize and train the models
models = {
    'LightGBM': lgb.LGBMClassifier(random_state=42, verbosity=-1),
    'Naive Bayes': GaussianNB(),
    'Support Vector Machine': SVC(probability=True, random_state=42),
    'AdaBoost': AdaBoostClassifier(n_estimators=100, random_state=42)
}

# Fit all models
for model_name, model in models.items():
    model.fit(X_train_preprocessed.toarray(), y_train)
    print(f"{model_name} Model Trained")

# SHAP Analysis
for model_name, model in models.items():
    print(f"\nSHAP Analysis for {model_name}")

    if model_name == 'LightGBM':
        explainer = shap.TreeExplainer(model)
        shap_values = explainer.shap_values(X_test_preprocessed.toarray())
    else:
        explainer = shap.KernelExplainer(model.predict_proba, X_train_preprocessed.toarray()[:100])
        shap_values = explainer.shap_values(X_test_preprocessed.toarray()[:100])

    shap.summary_plot(shap_values, X_test_preprocessed.toarray()[:100], feature_names=X_train.columns)
```

```
Collecting lime
  Downloading lime-0.2.0.1.tar.gz (275 kB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 275.7/275.7 kB 7.8 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from lime) (3.10.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from lime) (1.26.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from lime) (1.13.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from lime) (4.67.1)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.11/dist-packages (from lime) (1.6.1)
Requirement already satisfied: scikit-image>=0.12 in /usr/local/lib/python3.11/dist-packages (from lime) (0.25.2)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (3.4.2)
Requirement already satisfied: pillow>=10.1 in /usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (11.1.0)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (2.37.0)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (2025.2.18)
Requirement already satisfied: packaging>=21 in /usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (24.2)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (0.4)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.18->lime) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.18->lime) (3.5.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (4.56.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (1.4.8)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (3.2.1)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib->lime) (1.17.0)
Building wheels for collected packages: lime
  Building wheel for lime (setup.py) ... done
  Created wheel for lime: filename=lime-0.2.0.1-py3-none-any.whl size=283834 sha256=...
  Stored in directory: /root/.cache/pip/wheels/...
Successfully built lime
Installing collected packages: lime
Successfully installed lime-0.2.0.1
LightGBM Model Trained
Naive Bayes Model Trained
Support Vector Machine Model Trained
AdaBoost Model Trained
SHAP Analysis for LightGBM
```

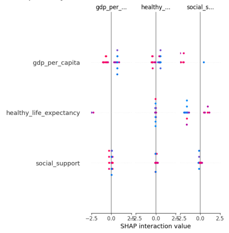

SHAP Analysis for Naive Bayes

100% 26/26 [01:24<00:00, 3.08s/it]

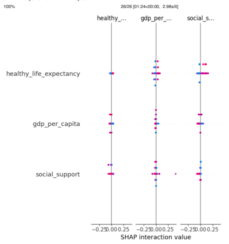

SHAP Analysis for Support Vector Machine

100% 26/26 [02:15<00:00, 5.16s/it]

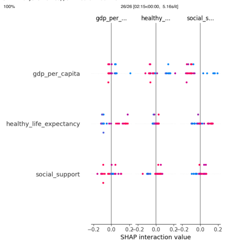

SHAP Analysis for AdaBoost

100% 26/26 [04:32<00:00, 10.85s/it]

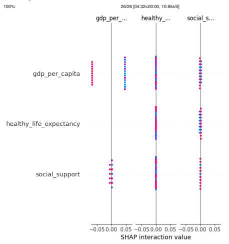


```python
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

# Extract feature names after preprocessing
# This code will get the full list of feature names including one-hot encoded columns
def get_feature_names(preprocessor, input_features):
    """Function to retrieve the full feature names including one-hot encoded columns."""
    output_features = []
    for name, transformer, columns in preprocessor.transformers_:
        if isinstance(transformer, Pipeline):
            transformer = transformer.steps[-1][1]
        if isinstance(transformer, OneHotEncoder):
            encoded_feature_names = transformer.get_feature_names_out(columns)
            output_features.extend(encoded_feature_names)
        else:
            output_features.extend(columns)
    return output_features
feature_names = get_feature_names(preprocessor, X_train.columns)

# LIME Analysis with Correct Feature Names
import lime
import lime.lime_tabular

for model_name, model in models.items():
    print(f"\nLIME Analysis for {model_name}")
    # Create a LIME explainer with correct feature names
    explainer = lime.lime_tabular.LimeTabularExplainer(
        X_train_preprocessed.toarray(),
        feature_names=feature_names,
        class_names=['Class 0', 'Class 1', 'Class 2'],
        discretize_continuous=True
    )
    idx = min(1, X_test_preprocessed.shape[0] - 1)
    exp = explainer.explain_instance(
        X_test_preprocessed.toarray()[idx],
        model.predict_proba,
        num_features=10
    )
    exp.show_in_notebook(show_all=False)
```

LIME Analysis for LightGBM

LIME Analysis for Naive Bayes

LIME Analysis for Support Vector Machine

LIME Analysis for AdaBoost

SHAP Analysis Interpretation

Across all models, gdp_per_capita, healthy_life_expectancy, and social_support are consistently among the top influential features. The LightGBM model shows a broader range of SHAP values, indicating a more dynamic response to feature changes, particularly with gdp_per_capita. In contrast, the AdaBoost model exhibits a smaller range of SHAP values, suggesting it is less sensitive to changes in input features compared to LightGBM.

The feature impact varies between models. In both LightGBM and SVM models, gdp_per_capita demonstrates a dual influence on predictions, with both positive and negative impacts. This dual behavior shows the feature's versatility in driving different outcomes depending on the context. The clustering of SHAP values near zero for some features in the Naive Bayes model suggests lower feature influence or a more linear response to those features. On the other hand, LightGBM and SVM models present a wider distribution of SHAP values, highlighting higher feature interaction and model complexity.

LIME Analysis Interpretation

LIME provides a local explanation for individual predictions, complementing the global insights from SHAP. For example, in the LightGBM model, life_expectancy and freedom_to_make_life_choices played significant roles in a specific prediction. The Naive Bayes model, however, focused primarily on categorical features (countries and regions), indicating it might rely more heavily on categorical groupings rather than nuanced numerical relationships.

LIME's visualizations for AdaBoost and SVM models reveal that predictions are influenced by both numerical and categorical features. However, the magnitude of influence in the AdaBoost model is generally smaller, which aligns with the narrower range of SHAP values observed earlier.

Overall Evaluation

SHAP offers a global view of feature importance across all predictions, effectively showing how each feature contributes to the model's decisions. This is particularly useful for understanding overall model behavior and ensuring alignment with domain knowledge. LIME, on the other hand, provides a local view for specific predictions, allowing for more precise justifications when explaining individual cases to